

ET0023 Operating Systems

5. Process Management & Scheduling

Process Management & Scheduling

- In a multiprogramming or multitasking system processes will be competing for
 - System resources
 - Memory
 - Files (I/O devices)
- OS has to manage these resources in a fair and efficient manner
- OS has to provide for Process Management & Scheduling

Process Creation

To create a new process, the OS has to

1. Name the process
2. Create a new Process ID and Process Control Block
3. Locate the program to be executed on disk and allocate memory for the code segment in RAM
4. Load the program into code segment, initialize PCB registers
5. Prioritize the process (default values)
6. Schedule the process for execution.

Lect05 Process Management & Scheduling

3


Process Creation

- Using CLI to create a process, new process becomes child of the command interpreter.
- Processes form a hierarchy and are linked by a tree structure.
- If Parent is signalled or killed, usually all Children received the same signal and are destroyed together with parent.
- When a Child is created, the process may
 - Duplicate the Parent process
 - Load a completely new program
- The Parent may also
 - Continue executing along side its Children
 - Wait for some/all to finish before proceeding

Lect05 Process Management & Scheduling

4

Process Hierarchies



Lect05 Process Management & Scheduling

5

Process Termination

- The process after creation, starts running and completes its job.
- The process must now terminate.
- Termination is usually one of the following
 1. Normal exit (voluntary)
 2. Error exit (voluntary)
 3. Fatal error (involuntary)
 4. Killed by another process (involuntary)

Lect05 Process Management & Scheduling

6

Process Scheduling

- To keep the CPU as busy as possible, the OS must perform scheduling.
- When a process blocks, it should be replaced.
- To let a process run for a maximum amount of time, an interrupt must be in place.
- Any stopped process must have its PCB saved and the next PCB restored.
- Process switching is expensive, hence the scheduler must be efficient.

Lect05 Process Management & Scheduling

7


CPU-I/O Burst Cycle

- All processes alternate bursts of computing (CPU cycles) with I/O requests.
- I/O wait is when a process enters the blocked state waiting for an external device to complete its work.
- CPU-bound: processes that use more time computing
- I/O-bound: processes that use more time waiting for I/O

Lect05 Process Management & Scheduling

8

How does a Process behave?



Bursts of CPU usage alternate with periods of IO wait


(a) A CPU-bound process

(b) An IO-bound process

Lect05 Process Management & Scheduling

9


Example



Lect05 Process Management & Scheduling

10

Modelling Multitasking



Lect05 Process Management & Scheduling

11


Analysis of Multi-Tasking Performance

Job	Arrival time	CPU minutes needed
1	10:00	4
2	10:10	3
3	10:15	2
4	10:20	2

(a)

	# Processes			
	1	2	3	4
CPU idle	.80	.64	.51	.41
CPU busy	.20	.36	.49	.59
CPU/process	.20	.18	.16	.15

(b)



(c)

Lect05 Process Management & Scheduling

12

Process Scheduling Policy

- An algorithm the Policy Scheduler uses to determine
 - How much CPU time does a process get
 - Who gets the process next (Priority)
 - When the process will be interrupted

Lect05 Process Management & Scheduling

13

Scheduling Criterion


- CPU Utilization: Objective to keep CPU as busy as possible
- Throughput: How many processes are completed per time unit
- Turnaround time: How long it takes to execute a process (from submission to completion)
- Waiting time: The time a process has to wait before execution starts
- Response time: Time from submission to obtaining the first results

Lect05 Process Management & Scheduling

14

Example

Four processes are submitted at time 0



P1	Waiting time	0
	Turnaround time	3
P2	Waiting time	$3 + 3 = 6$
	Turnaround time	9
P3	Waiting time	$4 + 4 = 8$
	Turnaround time	10
P4	Waiting time	5
	Turnaround time	7
Average waiting time		$(0+6+8+5)/4 = 4.75$
Average turnaround time		$(3+9+10+7)/4 = 7.25$

Lect05 Process Management & Scheduling

15

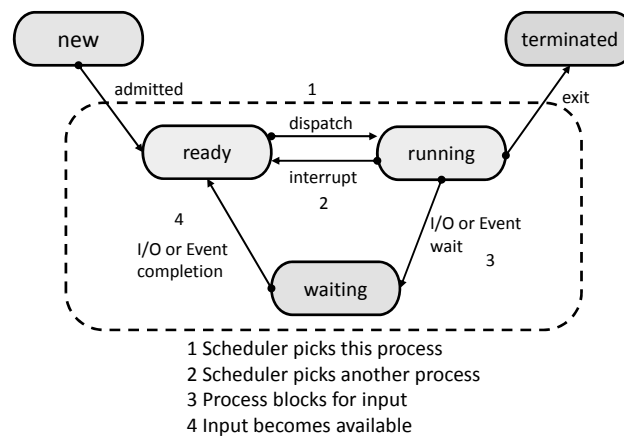
Process Scheduling Policy

- A good algorithm should
 - Maximize throughput
 - Minimize response time
 - Minimize turnaround time
 - Minimize waiting time
 - Maximize CPU efficiency
 - Ensure fairness for all processes

Lect05 Process Management & Scheduling

16

Process States



Lect05 Process Management & Scheduling

17

Scheduling Algorithms

- A non-preemptive scheduling algorithm picks a process to run and then just lets it run until it blocks (either on IO or waiting for another process) or until it voluntarily releases the CPU.
 - Does not require a clock.
- A preemptive scheduling algorithm picks a process and lets it run for a maximum of some fixed time. If it is still running at the end of time interval, it is suspended and the scheduler picks another process to run (if one is available).
 - Requires a clock.

Lect05 Process Management & Scheduling

18

Scheduling Environments

- Different environments require different scheduling algorithms
 - Batch → Background
 - Interactive → Foreground
 - Realtime → Multimedia
- There is no one size fits all!

Lect05 Process Management & Scheduling

19

Scheduling Goals

- All systems
 - Fairness – giving each process a fair share of the CPU
 - Policy enforcement – seeing that stated policy is carried out
 - Balance – keeping all parts of the system busy
- Batch systems
 - Throughput – maximize jobs per hour
 - Turnaround time – minimize time between submission and termination
 - CPU utilization – keep the CPU busy all the time
- Interactive systems
 - Response time – respond to requests quickly
 - Proportionality – meet user's expectations
- Real-time systems
 - Meeting deadlines – avoid losing data
 - Predictability – avoid quality degradation in multimedia systems

Lect05 Process Management & Scheduling

20

The “Perfect” Scheduler

- Minimize Latency
 - Response/Job completion time
- Maximize Throughput
 - Maximize job/time
- Maximize Utilization
 - Keep all devices busy
- Fairness
 - Every job makes progress, no job starves

Lect05 Process Management & Scheduling

21

Problem Analogy

- Imagine you are the cook at a “Chi Char” stall
 - Customers continually enter and place their orders
 - Dishes take varying amounts of time to prepare
- What is your goal?
- What strategy achieves your goal?


Lect05 Process Management & Scheduling

22

“Chi Char Stall” vs Multitasking

Chi Char Stall

- Waitress
- Cook
- Client
- Order
- Dish



Multi-tasking OS

- CPU Scheduler
- CPU
- User
- Process
- CPU or I/O Burst


Lect05 Process Management & Scheduling

23


CPU / I-O Burst

- A Process alternates between CPU and I/O Bursts

CPU bound: Long CPU bursts



I-O bound: Short CPU bursts




During an I/O burst = process is idle, switch to another “for free”

Lect05 Process Management & Scheduling

24

CPU Scheduler

- Processes and Threads migrate among queues (Ready, Device Wait, Run)
- Scheduler selects one from the ready queue to run
 - Which one?
 - When?




Lect05 Process Management & Scheduling

25

First Come First Served

- Jobs are scheduled in order of arrival
- Non-preemptive
- Problem: Average wait time depends on arrival order



- Advantage: Really Simple!!

Lect05 Process Management & Scheduling

26

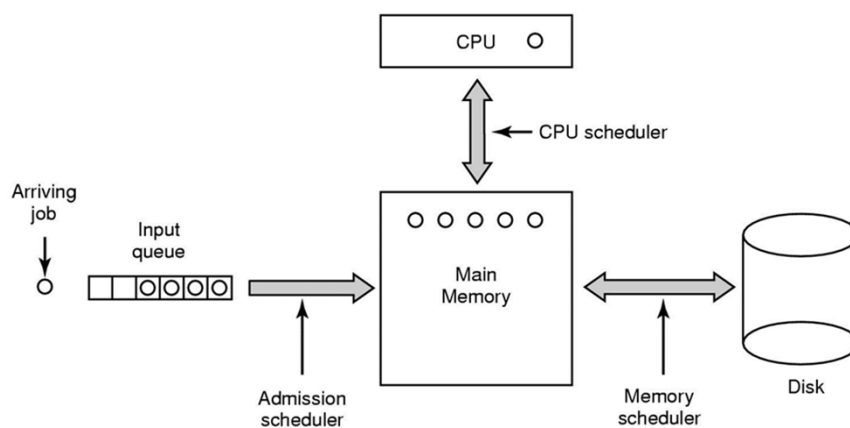
Batch Systems

- As jobs arrive at the system, they are initially placed in an input queue stored on the disk.
- The admission scheduler decides which jobs to admit to the system.
- The memory scheduler determines which processes are kept in memory and which on the disk.
- The CPU scheduler picks one of the ready processes in main memory to run next.

Lect05 Process Management & Scheduling

27

Scheduling in Batch Systems



Lect05 Process Management & Scheduling

28

Last-In-First-Out ?

- Newly arrived jobs are placed at the head of ready queue.
- Improves response time for newly created jobs
- Problem:
 - May lead to starvation (early processes may never get any CPU time!)

Lect05 Process Management & Scheduling

29

Round Robin

- FCFS with preemption
 - Often used for timesharing
 - Ready queue is treated as a circular queue (FIFO)
 - Each process is given a time slice (quantum)
 - Process runs for the quantum or until blocked
 - RR allocates the CPU uniformly (fairly) across participants
 - If average queue length is n , each quantum is $1/n$


Lect05 Process Management & Scheduling

30

Round Robin

Process Burst
Time (sec)
P1 53
P2 17
P3 68
P4 24

Time Quantum = 20



Lect05 Process Management & Scheduling

31

RR: Choice of Time Quantum


- Performance depends on length of time-slice
 - Context switching is not free
 - If time-slice is set too high
 - FCFS
 - Process will finish or block before their slice is up
 - If time-slice is set too low
 - Overhead is on the context switching between processes
- Usual Practice
 - Context switch is usually negligible (<1% per timeslice)
 - Context switch too often – lose all productivity

Lect05 Process Management & Scheduling

32

Shortest Job First

- Choose the job with the shortest next CPU burst.
- Optimal for minimizing average meeting time
- Problem: Impossible to know the length of next CPU burst.




Lect05 Process Management & Scheduling

33

Preemptive Shortest Job First

- SJF can be preemptive or non-preemptive
- When new job arrives and current job has long time to execute.
- Preemptive SJF = Shortest Remaining Time First (SRTF)



Lect05 Process Management & Scheduling

34

Priority Scheduling

- Choose next job based on priority
- For SJF, priority = expected CPU burst
- Can be preemptive or non-preemptive
- Can approximate to other scheduling algorithms
 - P (arrival time) FIFO
 - P (now – arrival time) LIFO
 - P (job length) SJF
- Problem:
 - Starvation: jobs can wait indefinitely
 - Solution: Age processes, increase priority as a function of waiting time.

Lect05 Process Management & Scheduling

35

Priority Scheduling

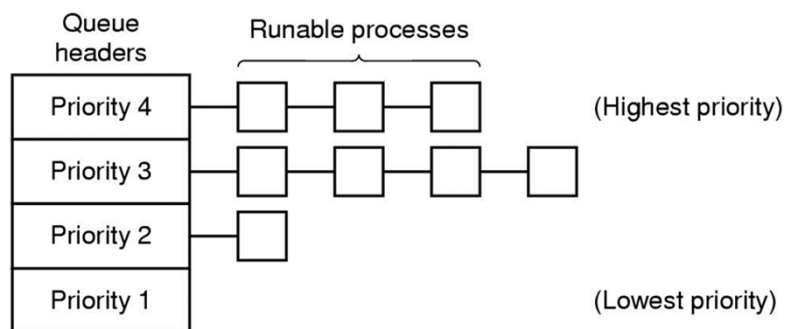
- Can use external characteristics of the process to assign priorities
 - Position of user (e.g. root/sysadmin)
 - Cost (\$\$\$ - he who pays gets first)
- Can use internal characteristics of the process to assign priorities
 - Memory requirements
 - Number and type of peripheral devices
 - Estimated CPU time
 - Amount of time already spent in the system

Lect05 Process Management & Scheduling

36

Priority Scheduling using Multiple Queues

- Each process is assigned a priority, and the runnable process with the highest priority is allowed to run.
- Idea used by Windows and Linux.

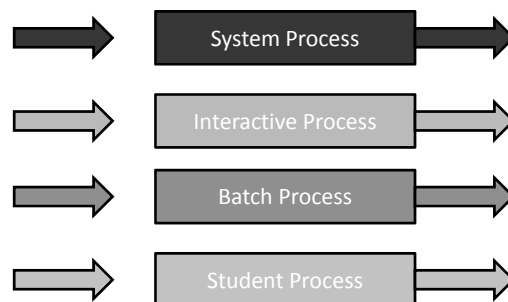


Lect05 Process Management & Scheduling

37

Multilevel Queue Scheduling

Highest Priority



Lowest Priority

Lect05 Process Management & Scheduling

38

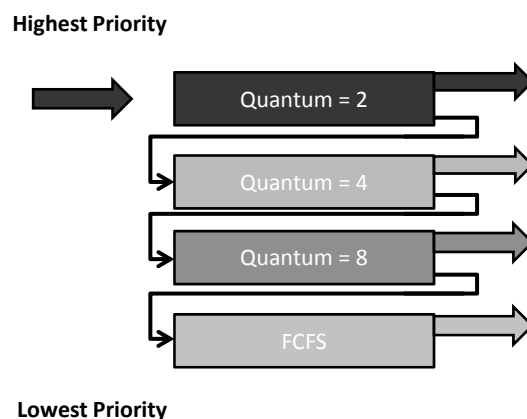
Multilevel Queue Scheduling

- Implement multiple ready queues based on job "type"
 - System processes
 - Interactive processes
 - Student programs
 - CPU-bound processes
 - Batch jobs
- Different queues may be scheduled using different algorithms
- Intra-queue CPU allocation is either **strict** or **proportional**
- Problem: classifying jobs into queues is difficult.

Lect05 Process Management & Scheduling

39

Multilevel Feedback Queues



Lect05 Process Management & Scheduling

40


Multilevel Feedback Queue

- Implement multiple ready queues
 - Different queues may be scheduled using different algorithms
 - Similar to multilevel queue scheduling but assignments are not static
- Jobs move from queue to queue based on feedback
 - Feedback = Behaviour of the job
 - e.g. full quantum for computation, or frequent I/O
- Very general algorithm, select parameters for
 - Number of queues
 - Scheduling algorithm for each queue
 - When to upgrade or downgrade a job

Lect05 Process Management & Scheduling

41

A Multi-level System



Lect05 Process Management & Scheduling

42

Real-time Scheduling

- Real-time processes have timing constraints, Expressed as
 - Deadlines
 - Rate requirements
- Common misconception:
 - Real time does not mean fast.

Lect05 Process Management & Scheduling

43

Real-time Scheduling

- Common RT scheduling policies
 - Rate monotonic
 - Just one scalar priority related to productivity of job
 - $\text{Priority} = 1/\text{rate}$
 - Static
 - Earliest Deadline First (EDF)
 - Dynamic but more complex
 - $\text{Priority} = \text{Deadline}$
- Both require admission control to provide guarantees

Lect05 Process Management & Scheduling

44

Real-Time System Scheduling

- Hard real-time
 - There are absolute deadlines that must be met, or else!
- Soft real-time
 - Missing an occasional deadline is undesirable, but nevertheless tolerable.

Lect05 Process Management & Scheduling

45

Thread Scheduling

- All thread share code & data segment
- Option 1: Ignore this fact
- Option 2: Two-level scheduling
 - User-level scheduling
 - Schedule process, and within each process, schedule threads
 - Reduce control switching overhead and improve cache hit ratio

Lect05 Process Management & Scheduling

46

Multiprocessor Scheduling

- Option 1: Ignore this fact
- Option 2: Space based affinity
 - Assign threads to processes
 - Control resource sharing
- Option 3: Gang scheduling
 - Run all threads belonging to a process at the same time
 - Low-latency communication
 - Greater distance (scheduling)

Lect05 Process Management & Scheduling

47

Problem Cases

- Blindness about job types
 - I/O goes idle
- Optimization involves favouring jobs of type "A" over "B"
 - Lots of "A"s, then "B" starve
- Interactive process trapped behind others
 - Response time suffers for no reason
- Priorities: A depends on B, A's priority > B's
 - B's never run

Lect05 Process Management & Scheduling

48

Convoy Effect

- CPU bound job will hold CPU until done, or it causes an I/O burst (rare)
 - Long periods where no I/O requests are issued, CPU is held
 - Result: Poor I/O device utilization
- Example: One CPU bound, many I/O bound
 - CPU-bound runs, I/O devices idle
 - CPU-bound blocks
 - I/O bound job(s) run, quickly blocks on I/O
 - CPU-bound runs again
 - I/O completes
 - CPU-bound still runs while I/O devices idle

Lect05 Process Management & Scheduling

49

- QUESTIONS

Lect05 Process Management & Scheduling

50