


ET0023 Operating Systems

4. Threads

Threads



http://en.wikipedia.org/wiki/Thread_%28computer_science%29

Processes vs Threads

- Processes
 - Expensive (in terms of time) to switch
 - Requires file, resource, accounting structures
 - Require a fully Process Control Block
 - Heavy weight
- Threads
 - Scheduled in a user-dependent way
 - Independent sub-processes that share the data/address space of the process
 - Requires on stack and registers
 - Light weight

Threads

- Mini-versions of a process.
- Enables the user to split up a process into logically separate parts (threads)
- Each thread is run separately and independently of one another until they need to communicate.
- Allow certain functions to be executed in parallel with others.

Threads

- Multi-threaded tasks are comprised of several independent processes which work on common data, and can run in parallel.
- Threads only require stack and some registers. All other resources are shared.
- Lightweight
- Can be assigned priorities
- User controlled in the program.

Thread Model

Per process items

Address space
Global variables
Open files
Child processes
Pending alarms
Signals and signal handlers
Accounting information


Items shared by all
threads in a process

Per thread items

Program counter
Registers
Stack
State

Items private to
each thread

The Thread Model



Each thread has its own

- Thread ID
- Program Counter
- Register Set
- Stack Space

Why use threads?

- If the scheduler knows what resources a process requires, then the sharing of resources could be made more efficient.
(Scheduler is a program – programs cannot think!!)
- If there are multiple CPUs, you could run multiple processes in parallel – but what about resources ?
- When you switch processes, you have to save the whole PCB and restore the next PCB (slow, CPU intensive)
- What if we could "think" for the scheduler?

Advantages of threads


- Allows a programmer to switch between lightweight processes when it is best for the program
- Each process gets a slice of the CPU time, but what if we can make parts of the process work in parallel
- Responsiveness: threads increase responsiveness to a process
- Resource sharing: threads share the resources
- Economy: thread CPU cost is low
- Multiprocessor architectures allow multiple threads to run.

http://en.wikipedia.org/wiki/Multithreading_%28computer_architecture%29

Use of Threads

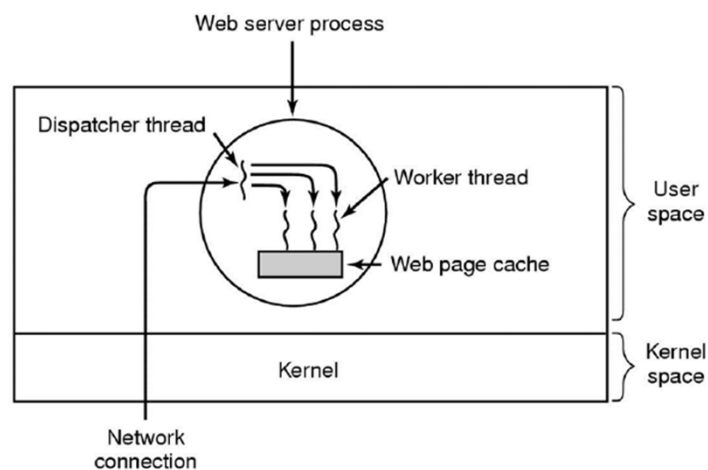
- Interactive applications.
 - Helps make apps more interactive.
 - Eg. Word processor, Excel spreadsheet.
- Server applications.
 - Serve multiple requests at the same time.
 - Eg. Web server.
- Application that process large amounts of data.
 - Input thread, Processing thread, Output thread.

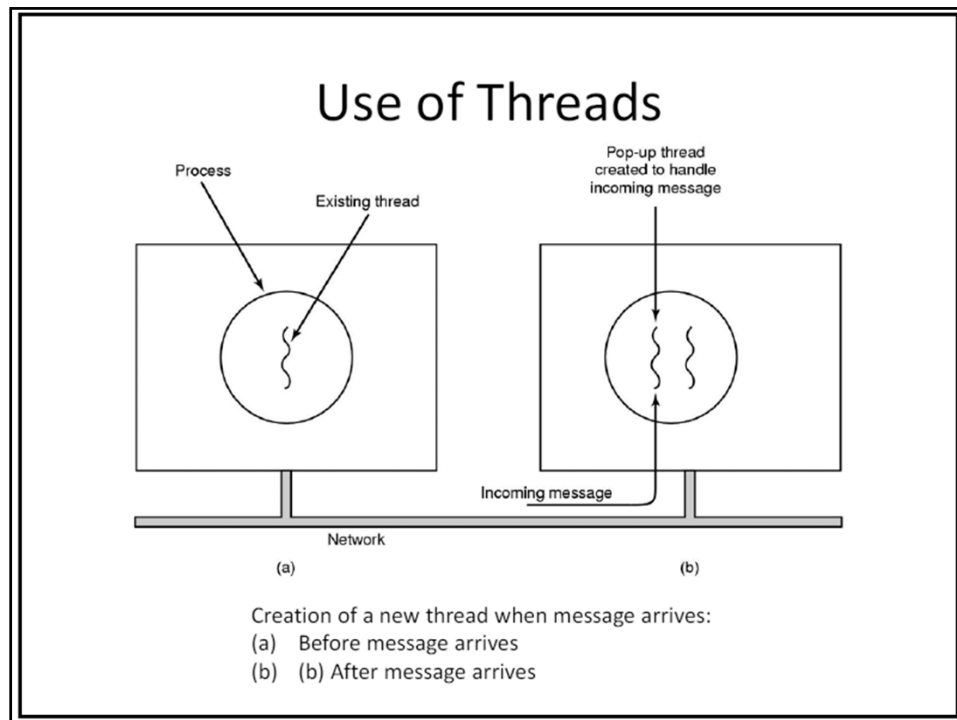
Threaded Word Processor



A word processor with three threads

Multi-threaded Web Server






Thread Types - User

- User Thread
 - Supported above the kernel
 - Implemented by a thread library at user level.
 - Library provides support for thread creation, scheduling and management.
 - No support from the kernel, no kernel intervention
 - Fast to create and manage
 - Problem if one thread performs a blocking system call, all threads will be blocked.


User-level threads



Thread Types - Kernel

- Kernel Thread
 - Supported directly by kernel (operating system)
 - Kernel performs thread creation, scheduling and management in kernel space.
 - Slower to create and manage
 - In multiprocessor environments, kernel can schedule threads on different processors.
 - Blocking calls are handled by kernel, the kernel can schedule another thread in the application.

Kernel-level Threads



QUESTIONS?