

# ET0023 Operating Systems

## 6. Inter-process Communication

### Synchronization

- How can I synchronize the execution of multiple threads of the same process?
  - Problem Definition
    - Race condition
    - Critical-selection problem
  - Solutions
    - Spinlocks
    - Semaphores
  - Usage

## Problem Context

- Multiple threads of the same process have:
  - Private registers and stack memory (the context switching mechanism saves and restores registers when switching from thread to thread)
  - Shared access to the remainder of the process “state”
- Preemptive CPU Scheduling:
  - The execution of a thread is interrupted unexpectedly.
- Multiple cores executing multiple threads of the same process.

## Share Counting



Uncle Scrooge is a Disney character that loved money

- Uncle Scrooge wants to count his \$1 coins.
- Initially, he uses one thread that increases a variable `coin_counter` for every \$1 coin
- To accelerate the counting, he uses 2 threads but maintains the same variable `coin_counter`, but shared between the threads.



## Share Counting

```
coin_counter = 0
```

### Thread A

```
while (mc_A_has_coins)
    ++coins_counter ;
```

### Thread B

```
while (mc_B_has_coins)
    ++coins_counter ;
```

```
print coins_counter
```

- What might go wrong?



## Share Counting

### Thread A

```
r1 = coins_counter
r1 = r1 + 1
coins_counter = r1
```

### Thread B

```
r2 = coins_counter
r2 = r2 + 1
coins_counter = r2
```

- If `coins_counter = 20`, what are the possible values after the execution of one A/B loop?

## Shared Counters

- One possible result: everything works!
- Another possible result: lost update!  
Difficult to debug.
- Called a "race condition"

## Race Conditions

- Def: A timing dependent error involving shared state
- Depends on how threads scheduled
  - Thread A starts
  - Thread A needs to "race" to finish it, because
  - Thread B looks at shared area and changes it
  - Thread A's change will be lost.
- Hard to detect:
  - All possible schedules have to be safe
    - Number of possible schedule permutations is huge
    - Some bad schedules? Some that will work sometimes?
  - Race conditions are intermittent
    - Timing dependent = small changes can hide bug



## Share Counting

```
coin_counter = 0
```

### Thread A

```
while (my_mc_has_coins)
```

```
    Enter critical-selection
    coins_counter ++
    Exit critical-selection
```

### Thread B

```
while (my_mc_has_coins)
```

```
    Enter critical-selection
    coins_counter ++
    Exit critical-selection
```

```
print coins_counter
```

## Critical Selection Problem

- The solution should satisfy
  - Mutual exclusion
  - Progress
  - Bounded waiting

```
Enter critical-section
```

```
    Critical section commands
```

```
Exit critical-section
```

```
    Remainder section
```

### Thread A

```
Enter section
    r1 = coins_counter
    r1 = r1 + 1
    coins_counter = r1
Exit section
```

### Thread B

```
Enter section
    r2 = coins_counter
    r2 = r2 + 1
    coins_counter = r2
Exit section
```

## Solution



- LOCK
- A process must acquire a LOCK to enter a critical section.
- Can be implemented using hardware or software

## Test-And-Set

- CPU **Hardware** instruction
- Test and modify the content of one word automatically
- Spinlock

```
bool testAndSet(bool *target) {  
    bool rv = *target;  
    *target = TRUE;  
    return rv;  
}
```



## Share Counting

```
coin_counter = 0
```

```
lock = false
```

### Thread A

```
while (my_mc_has_coins)
  while (testAndSet(&lock)) {
    coins_counter++;
    lock = false;
  }
}
```

### Thread B

```
while (my_mc_has_coins)
  while (testAndSet(&lock)) {
    coins_counter++;
    lock = false;
  }
}
```

```
print coins_counter
```

## Swap

- CPU **Hardware** instruction
- Swap the contents of two **words**  
**Automatically**
- Spinlock

```
void swap (bool *a, bool *b) {
  bool temp = *a;
  *a = *b;
  *b = temp;
}
```



## Share Counting

```
coin_counter = 0
```

```
lock = false
```

### Thread A

```
while (my_mc_has_coins)
  keyA = true
  while (keyA == true)
    swap (&lock, &keyA);
    coins_counter ++;
    lock = false;
  }
```

### Thread B

```
while (my_mc_has_coins)
  keyB = true
  while (keyB == true)
    swap (&lock, &keyB);
    coins_counter ++;
    lock = false;
  }
```

```
print coins_counter
```

## Problem

- TestAndSwap and Swap
  - Do not satisfy the bounded-waiting requirement.
  - Are complicated for application programmers to use
- What is the alternative?



## Semaphores

- Integer values
- Atomic operations
  - wait()
    - It decrements the value of semaphore variable by 1.
    - If the value becomes negative the process executing wait() is blocked, i.e., added to the semaphore's queue.
  - signal()
    - It increments the value of semaphore variable by 1.
    - After the increment if the value is negative, it transfers a blocked process from the semaphore's queue to the ready queue.

```
wait(S) {
    while ( S <= 0 );
    S--;
}
```

```
signal (S) {
    S++;
}
```



## Share Counting

```
coin_counter = 0
```

```
S = 1
```

### Thread A

```
while (my_mc_has_coins)
    wait(S);
    coins_counter ++;
    signal(S);
}
```

### Thread B

```
while (my_mc_has_coins)
    wait(S);
    coins_counter ++;
    signal(S);
}
```

```
print coins_counter
```

## Spinlock

- This implementation of Semaphors, TestAndSet and Swap require busy waiting.
- The waiting processes should loop continuously in the entry code.
- Valuable CPU cycles are wasted.
- Solution:
  - Block the waiting process.
  - Signal blocked process when the semaphore is “available”

## Semaphores

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

```
wait (semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
```

```
signal (semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

## Usage

- Binary semaphore (mutex)
  - Ranges between 0 and 1
  - E.g. Only one process can access a resource
- Counting semaphore
  - Ranges between 0 and N
  - E.g. N resources and M processes that share the resources
- Synchronization
  - Ranges between 0 and 1
  - E.g. Process A should do task  $A_t$  after B having done task  $B_t$

- QUESTIONS